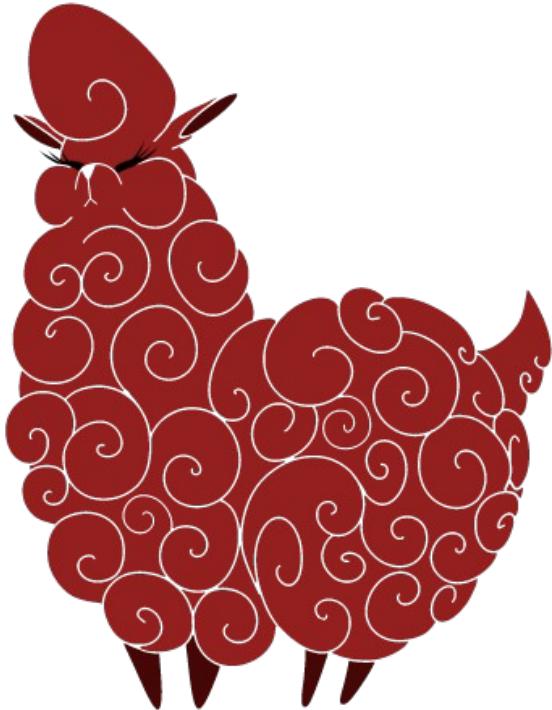


# Lisp "images" in LBM



Bo Joel Svensson  
A Chalmers FP-Talk!  
2025

# LispBM (LBM)

- Lisp like language for microcontrollers.
  - (,), car, cdr, (+ 1 2)
- With some built-in functionality inspired by Erlang.
  - Pattern matching, message passing, process monitoring.

# Target platforms

- Flexible, but mostly
  - STM32 (ARM) microcontrollers.
    - 192KB ram, up to 1MB flash.
  - ESP32C3 (RISC-V) microcontrollers.
    - 400KB ram, 4MB flash.
- Also runs fine on
  - X86, Raspberry PI ...

# Purpose

- Integrate into existing embedded application to provide a scripting layer.
  - Sandboxed.
  - Not the only thing using FLASH/RAM.

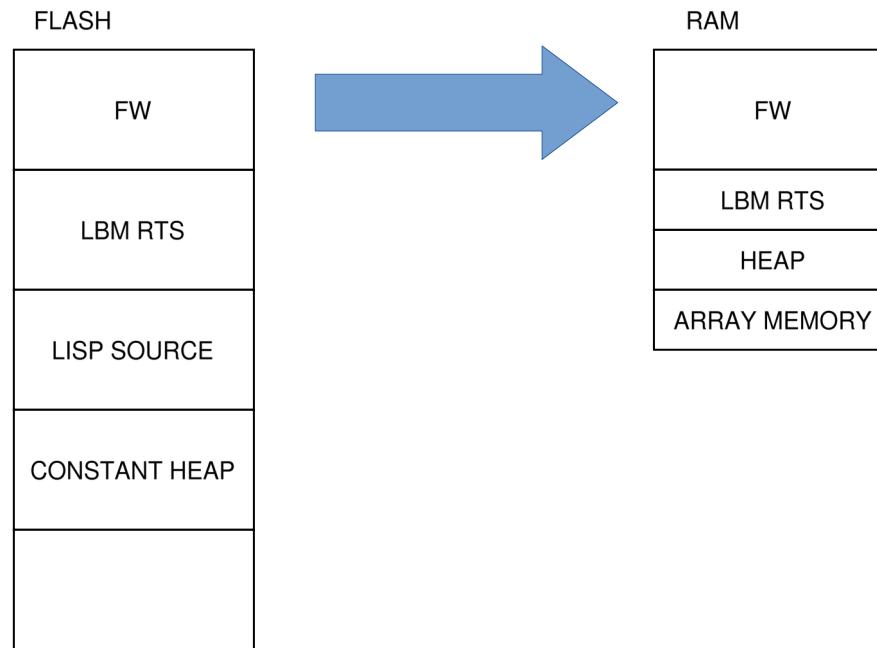
# Problems!

- Our programs are larger than our RAM.
- Flash memory has limitations.
- Booting is slow.

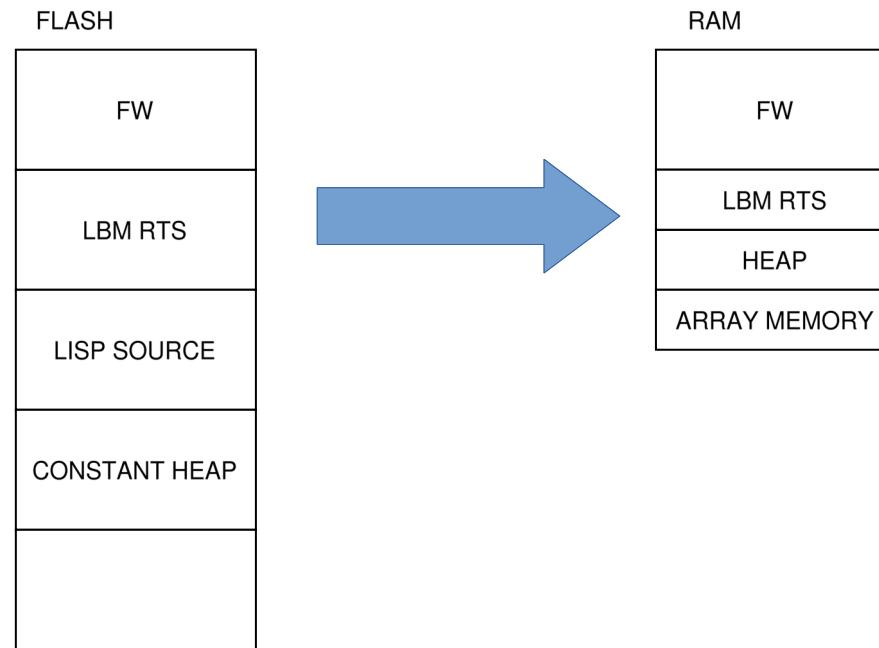
# Booting an LBM application



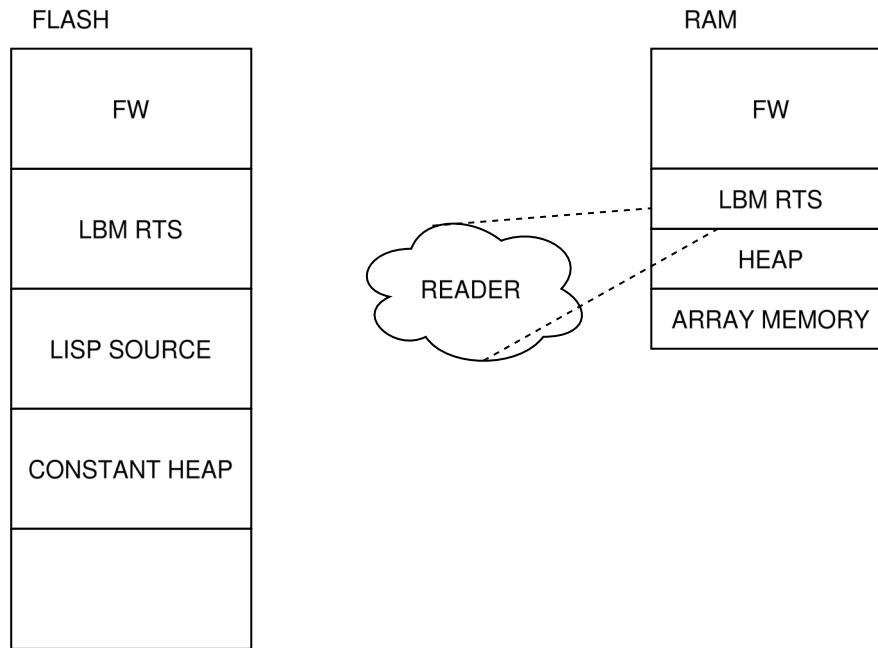
# Boot process



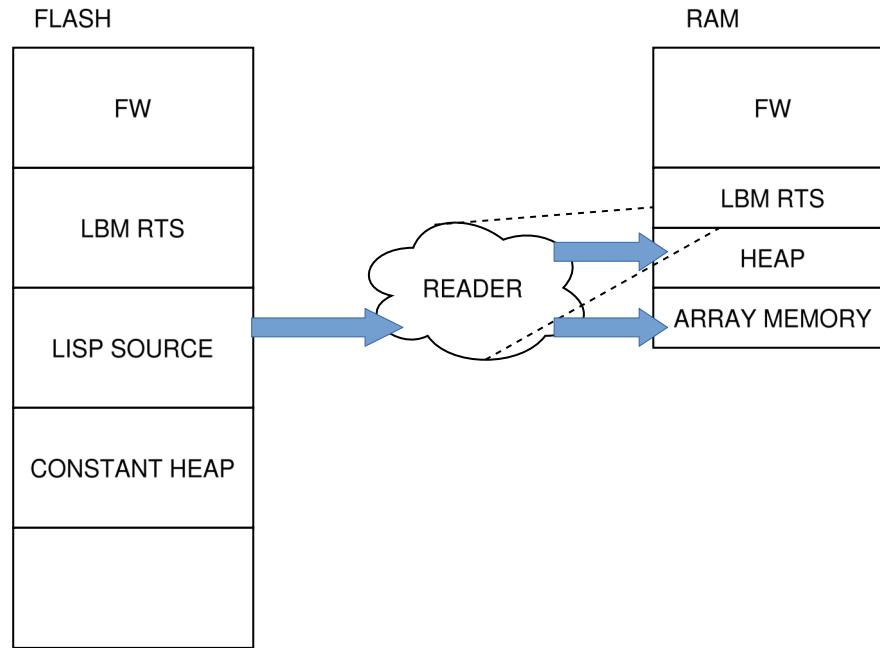
# Boot process



# Boot process



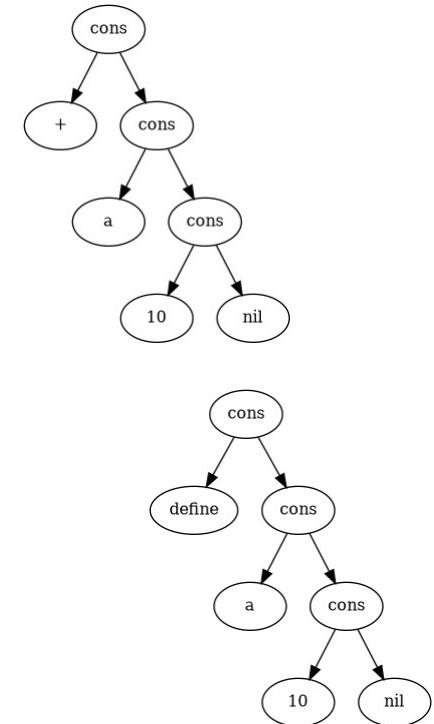
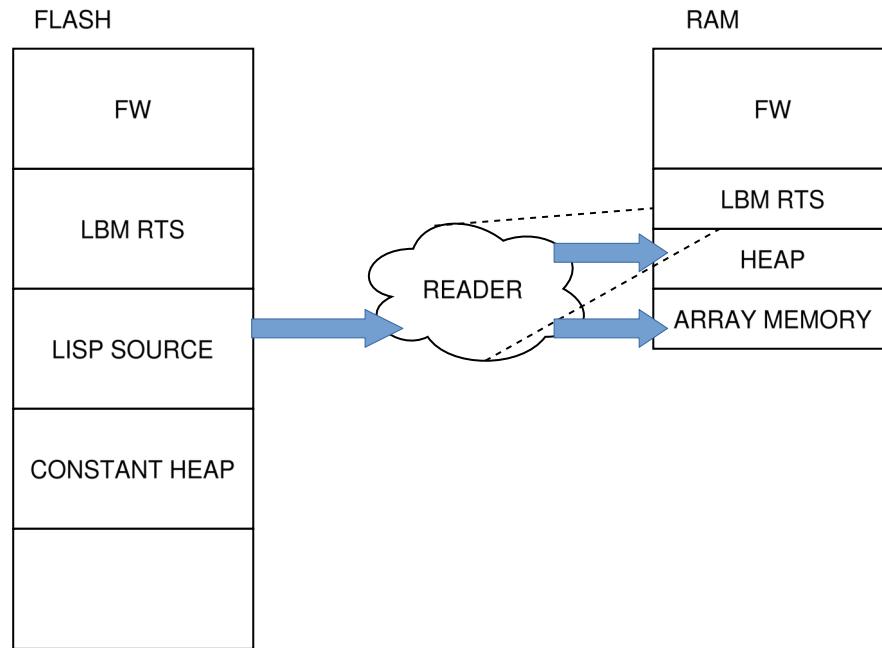
# Boot process



# Boot process

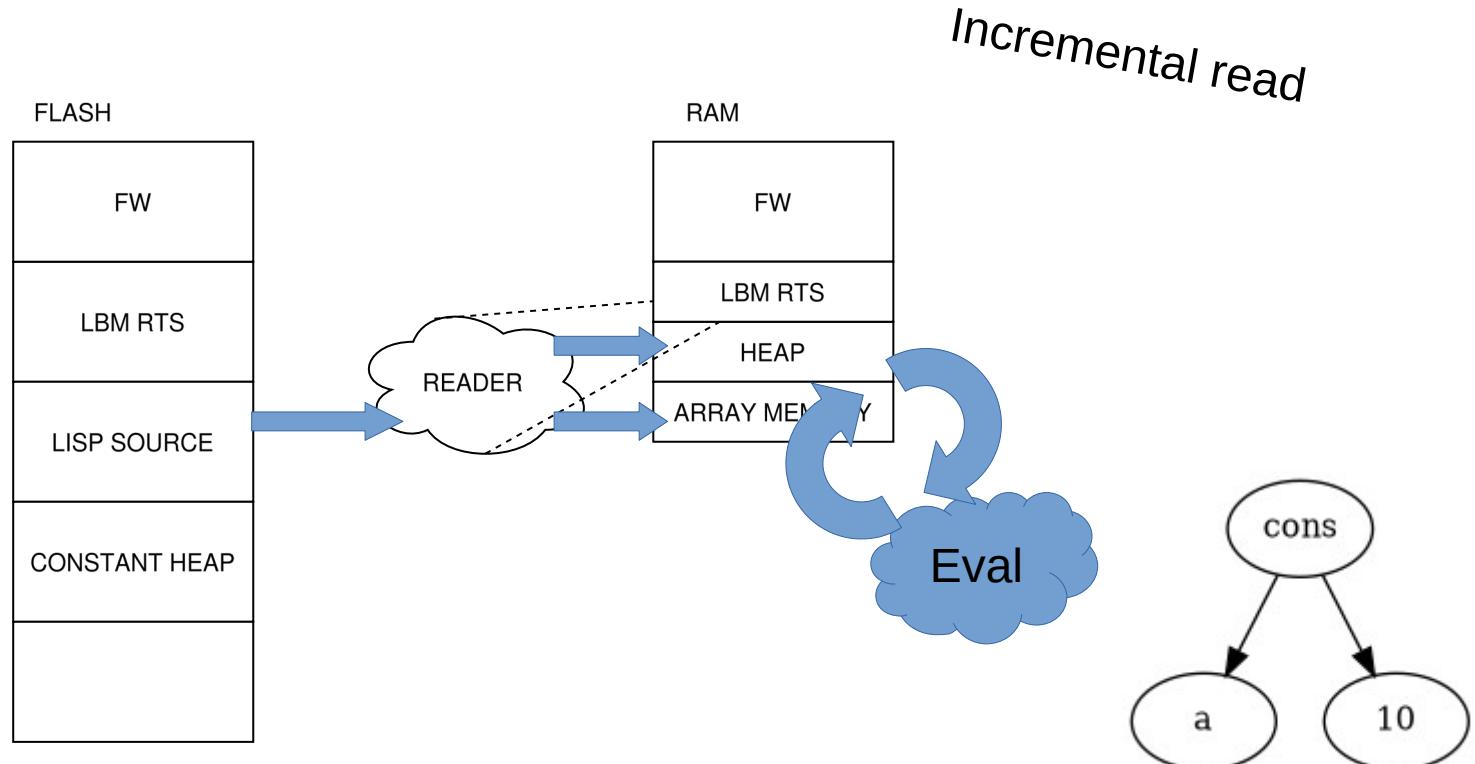
(define a 10)

(+ a 10)



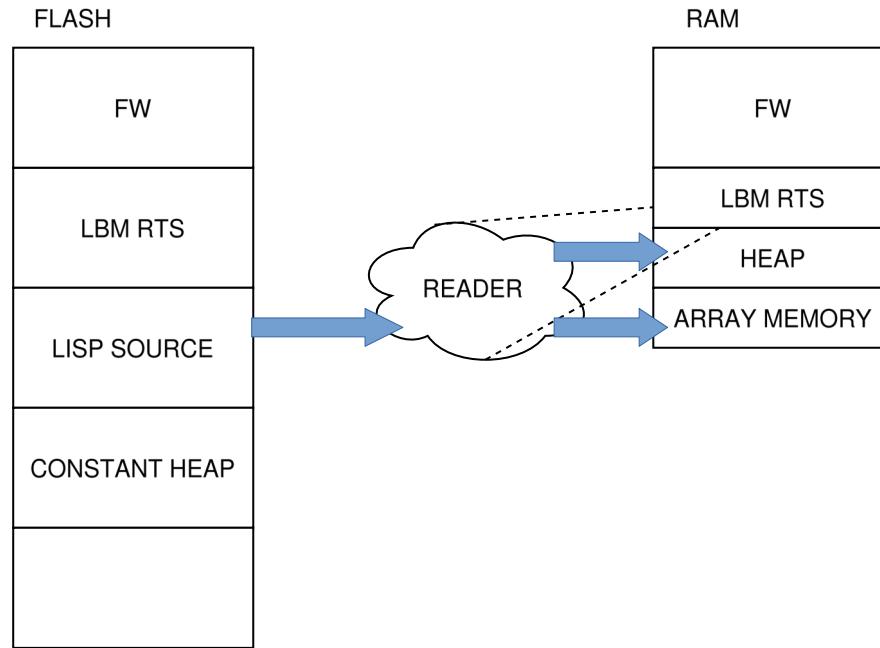
# Boot process

(define a 10)  
(+ a 10)

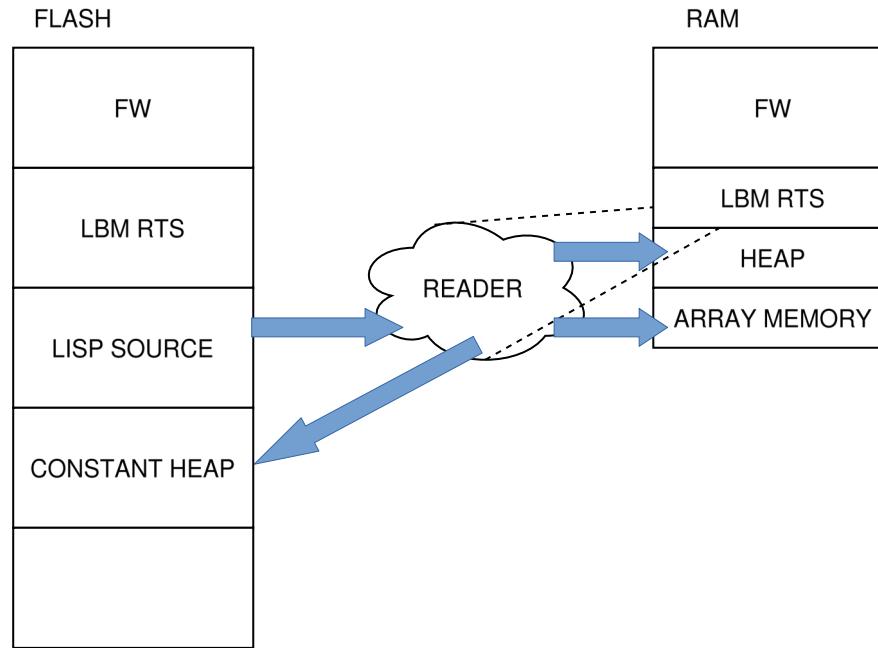


# Boot process

May require  
more heap than  
we have!

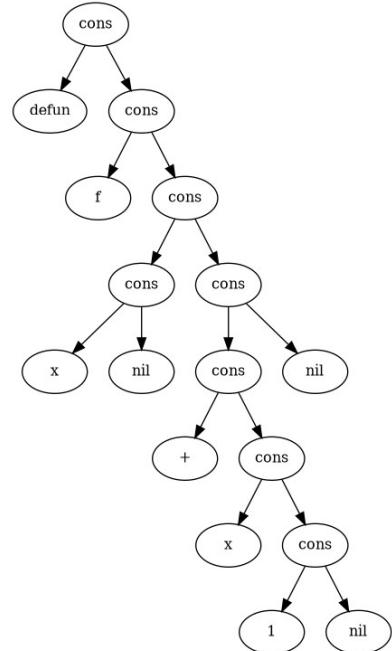
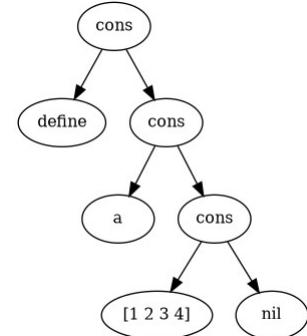
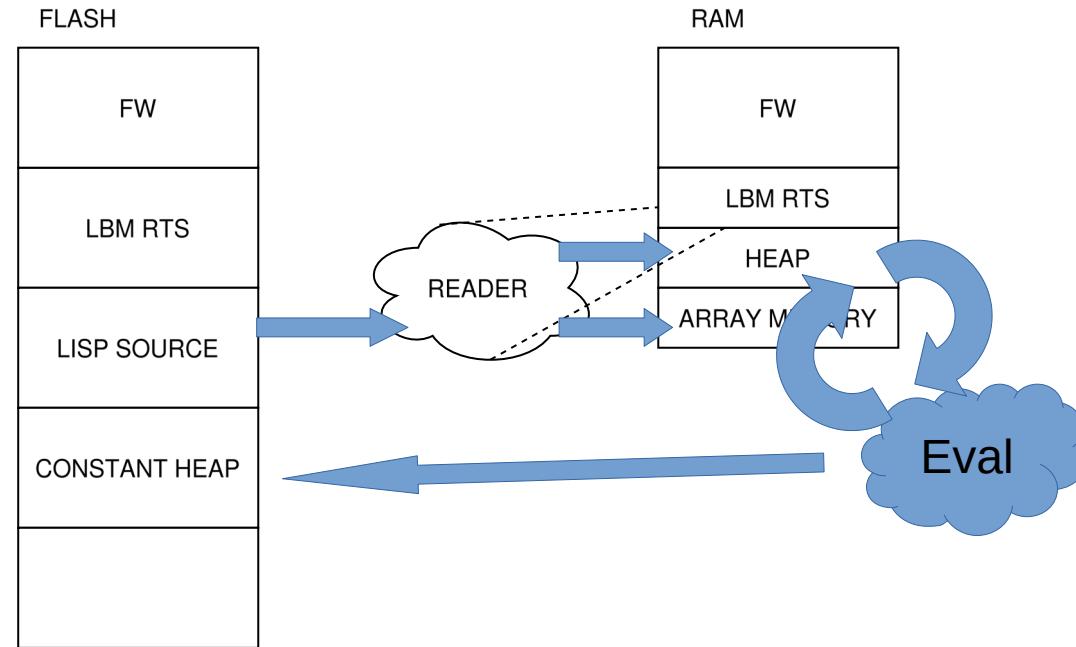
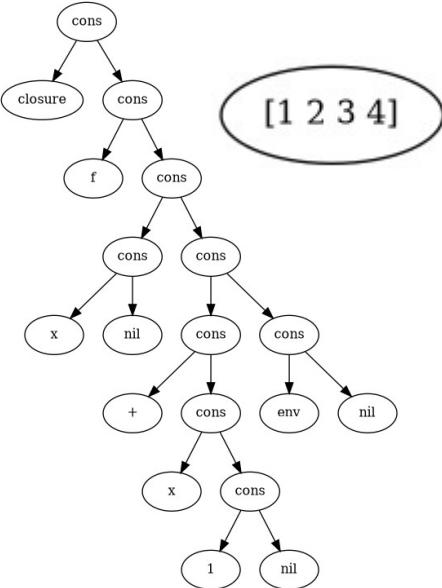


# Boot process



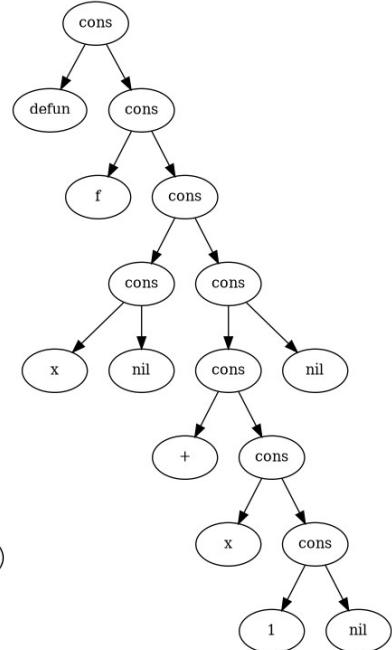
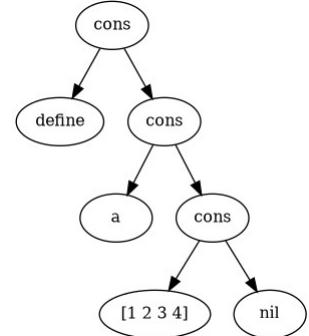
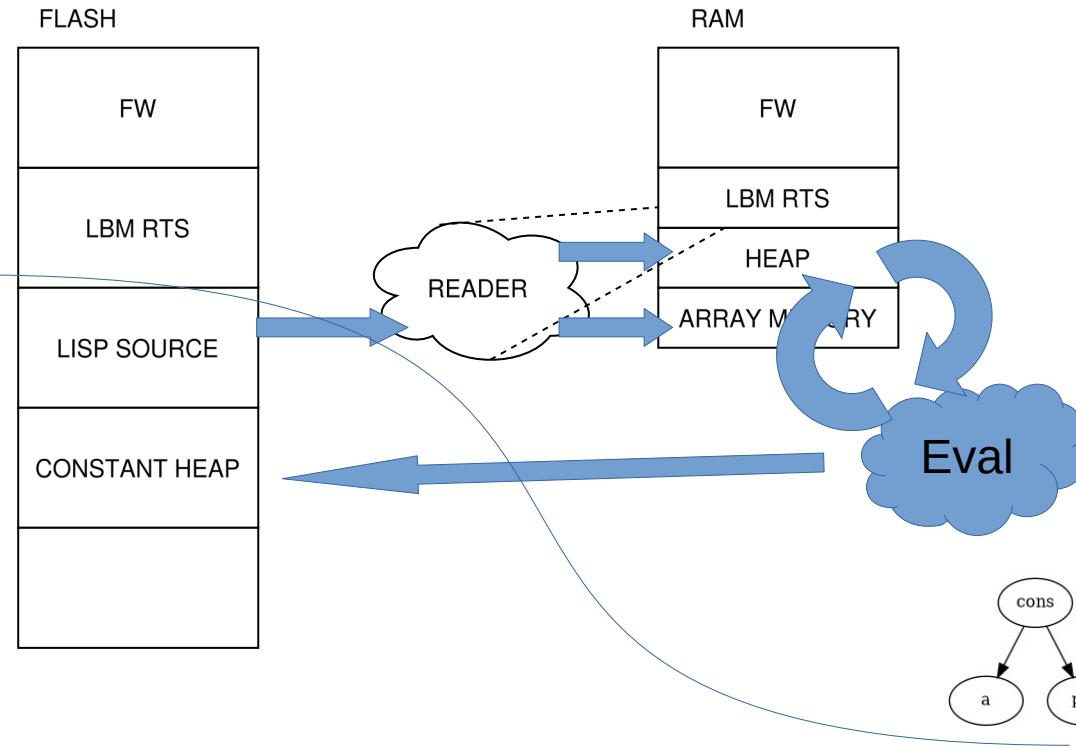
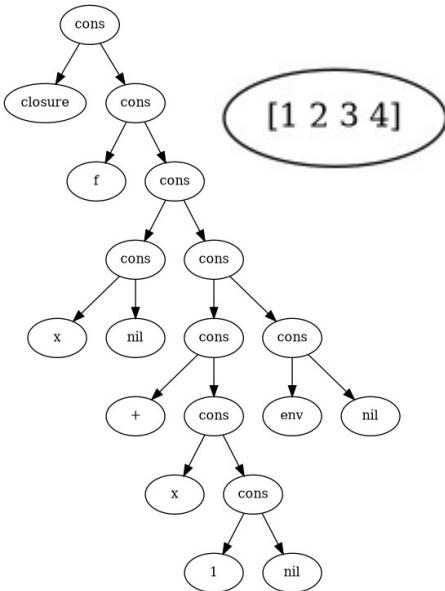
# Boot process

```
@const-start
(define a [1 2 3 4])
(defun f (x) (+ x 1))
@const-end
```



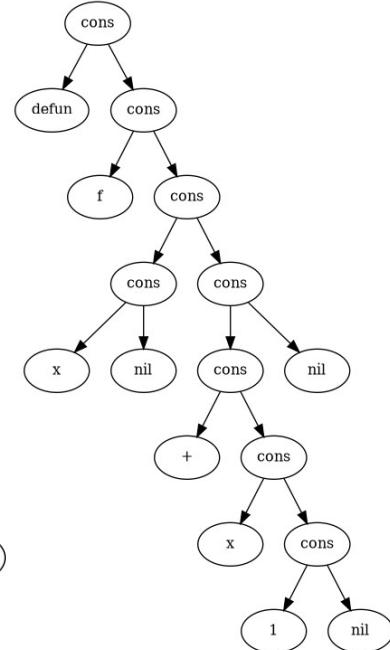
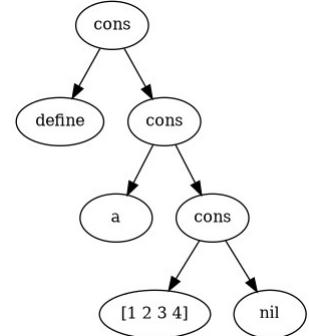
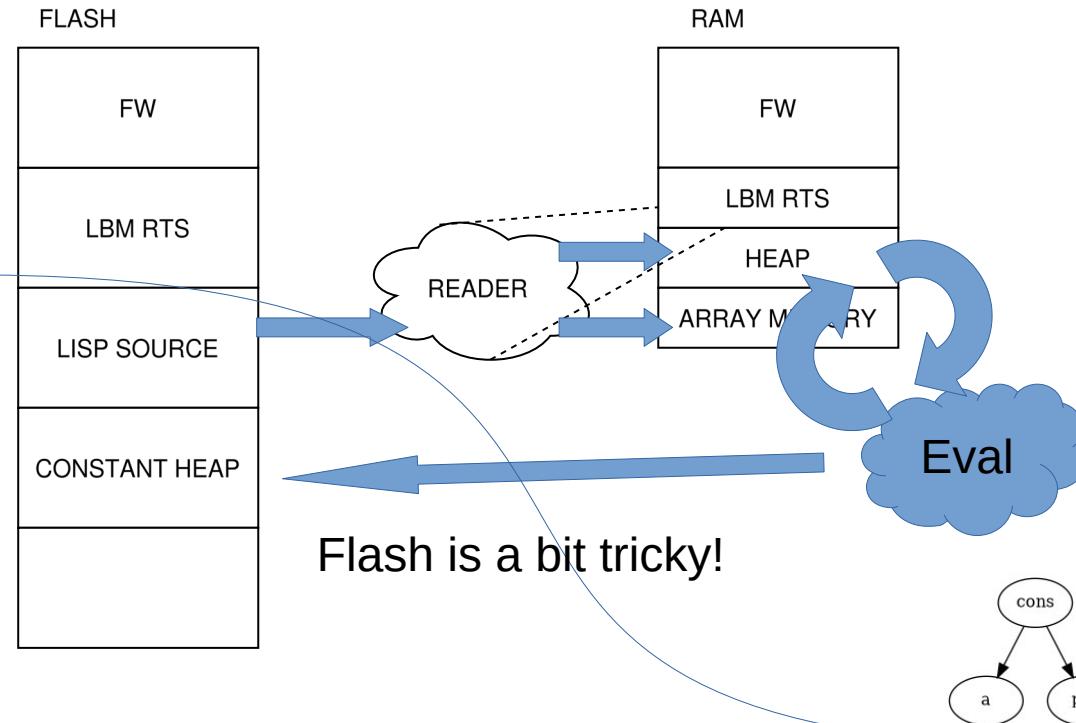
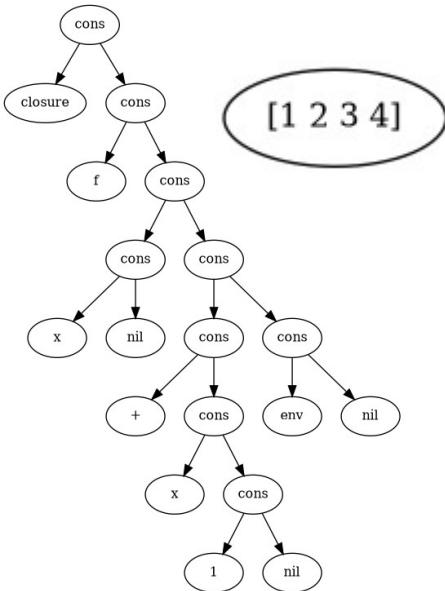
# Boot process

```
@const-start
(define a [1 2 3 4])
(defun f (x) (+ x 1))
@const-end
```



# Boot process

```
@const-start
(define a [1 2 3 4])
(defun f (x) (+ x 1))
@const-end
```



# Why are constant blocks tricky

```
(spawn f)

(if (system-initialized)
  ()
  (initialize))

(define randomly-long-list (range (random 10)))

@const-start
(stuff)
@const-end
```

# Summary so far

- The boot process is slow!
- Constant heap is a bit awkward.
  - Careful programming needed!

# Images to the rescue!

- Image based development.
  - Smalltalk, SBCL.
  - Fun playful REPL-interaction based development.
  - (save-image)

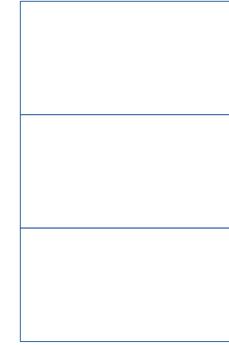
# The idea

Distributable  
binary blob

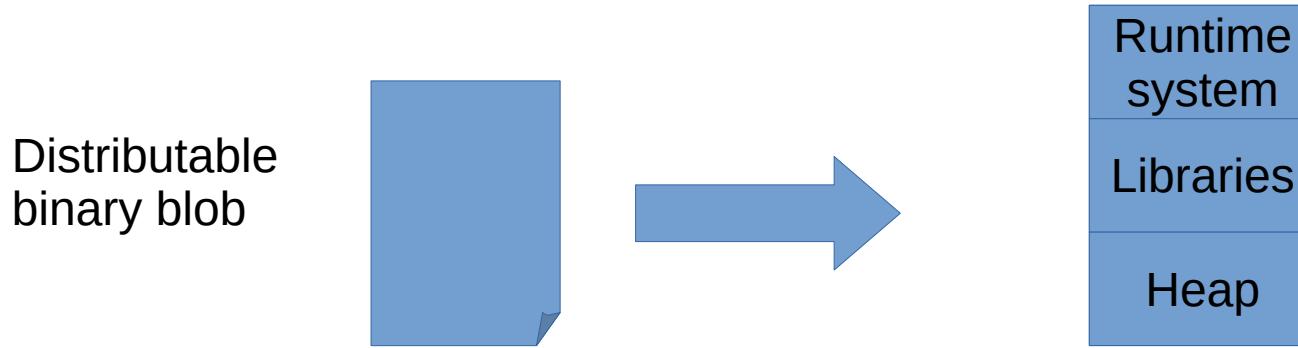


# Restart

Distributable  
binary blob



# Restart



Restarting a system in this way should be MUCH faster than that entire reader based bootup-process from earlier!



# Problems!

- Not enough room for a complete copy of HEAP, arrays-memory and constant-heap.
- RTS data-structures must be restored upon image-boot.
  - Symbol numbering, extensions
- Initialization of peripherals.

# Decisions

- Not going to save all state.
  - Not saving threads.
  - Not even going to try to automatically reinitialize any peripherals.
  - Not going to save entire RAM heap to image.
  - Not going to save entire arrays memory to image.
- A Startup entry is saved in the image.
  - Reinitialize peripherals.
  - Start up threads.
- Save the values stored in the global environment.

```
(defun f () ... )  
(defun g () ... )  
  
(defun main () {  
    (init-uart 115200)  
  
    (spawn f)  
    (spawn g)  
  
})  
  
(image-save)
```

# Decisions

- Create the image incrementally
  - The constant heap is built-up inside of the image from the beginning.
  - RTS data-structures that needs to be restored upon image-boot are created inside the image storage by default.
    - Symbol string  $\leftrightarrow$  number mappings.
    - Extensions function pointer addresses.
    - Constant heap write position.

# Image Structure

- The image is a collection of data fields written into flash.
- Duplicated data fields are allowed, where later fields have priority over earlier. Allows Incremental work towards same image.

```
#ifdef LBM64
#define IMAGE_INITIALIZED (uint32_t)0xBEEF4001
#else
#define IMAGE_INITIALIZED (uint32_t)0xBEEF2001
#endif

#define CONSTANT_HEAP_IX    (uint32_t)0x02
#define BINDING_CONST       (uint32_t)0x03
#define BINDING_FLAT        (uint32_t)0x04
#define SYMBOL_ENTRY         (uint32_t)0x06
#define SYMBOL_LINK_ENTRY   (uint32_t)0x07
#define EXTENSION_TABLE     (uint32_t)0x08
#define VERSION_ENTRY        (uint32_t)0x09
```

# The image



- Constant heap grows from bottom of image and upwards.
- All other data is added from the top and downwards in the image.
- When these two write pointers meet, image is full.

# Creating an image

```
@const-start
(define a ...)
(define b ...)

(defun h ...)
@const-end

(define i ...)
(defun f () ... )
(defun g () ... )

(defun main ()
  ...
 ))

(image-save)
```

# Creating an image

Source is read and evaluated, data-structures built in constant heap.

```
@const-start
(define a ...)
(define b ...)

(defun h ...)
@const-end

(define i ...)
(defun f () ... )
(defun g () ... )

(defun main () {
    ...
})

(image-save)
```

# Creating an image

Source is read and evaluated, data-structures built in RAM heap.

```
@const-start
(define a ...)
(define b ...)

(defun h ...)
@const-end

(define i ...)
(defun f () ... )
(defun g () ... )

(defun main () {
    ...
})

(image-save)
```

# Creating an image

Saves RAM heap structures to image (and more)

```
@const-start
(define a ...)
(define b ...)

(defun h ...)
@const-end

(define i ...)
(defun f () ... )
(defun g () ... )

(defun main () {
    ...
})
(image-save)
```

# Typical Image-save

```
lbtm_value ext_image_save(lbtm_value *args, lbtm_uint argn) {
    (void) args;
    (void) argn;

    bool r = lbtm_image_save_global_env();

    lbtm_uint main_sym = ENC_SYM NIL;
    if (!lbtm_get_symbol_by_name("main", &main_sym)) {
        lbtm_value binding;
        if (!lbtm_global_env_lookup(&binding, lbtm_enc_sym(main_sym))) {
            if (lbtm_is_cons(binding) && lbtm_car(binding) == ENC_SYM_CLOSURE) {
                goto image_has_main;
            }
        }
    }
    lbtm_set_error_reason("No main function in image\n");
    return ENC_SYM_ERROR;
image_has_main:
    r = r && lbtm_image_save_extensions();
    r = r && lbtm_image_save_constant_heap_ix();
    return r ? ENC_SYM_TRUE : ENC_SYM NIL;
}
```

```
bool lbm_image_save_global_env(void) {
    lbm_value *env = lbm_get_global_env();
    if (env) {
        for (int i = 0; i < GLOBAL_ENV_ROOTS; i++) {
            lbm_value curr = env[i];
            while(lbm_is_cons(curr)) {
                lbm_value name_field = lbm_caar(curr);
                lbm_value val_field = lbm_cdr(lbm_car(curr));

                if (lbm_is_constant(val_field)) {
                    write_u32(BINDING_CONST, &write_index, DOWNWARDS);
                    write_lbm_value(name_field, &write_index, DOWNWARDS);
                    write_lbm_value(val_field, &write_index, DOWNWARDS);
                } else {
                    int fv_size = image_flatten_size(val_field);
                    if (fv_size > 0) {
                        fv_size = (fv_size % 4 == 0) ? (fv_size / 4) : (fv_size / 4) + 1; // num 32bit words
                        if ((write_index - fv_size) <= (int32_t)image_const_heap.next) {
                            return false;
                        }
                        write_u32(BINDING_FLAT, &write_index, DOWNWARDS);
                        write_u32((uint32_t)fv_size, &write_index, DOWNWARDS);
                        write_lbm_value(name_field, &write_index, DOWNWARDS);
                        write_index = write_index - fv_size; // subtract fv_size
                        if (image_flatten_value(val_field)) { // adds fv_size backq
                            // TODO: What error handling makes sense?
                            fv_write_flush();
                        }
                        write_index = write_index - fv_size - 1; // subtract fv_size
                    } else {
                        return false;
                    }
                }
                curr = lbm_cdr(curr);
            }
        }
        return true;
    }
    return false;
}
```

# Boot

- Check if there is an image: 0xBEEF2001
  - Constant heap exists. Initialize
  - Start reading fields from top of image.

# Restore environment

```
case BINDING_CONST: {
    lbm_uint bind_key = read_u32(pos);
    lbm_uint bind_val = read_u32(pos-1);
    pos -= 2;
    lbm_uint ix_key = lbm_dec_sym(bind_key) & GLOBAL_ENV_MASK;
    lbm_value *global_env = lbm_get_global_env();
    lbm_uint orig_env = global_env[ix_key];
    lbm_value new_env = lbm_env_set(orig_env,bind_key,bind_val);

    if (lbm_is_symbol(new_env)) {
        return false;
    }
    global_env[ix_key] = new_env;
} break;
```

# Restore environment

```
case BINDING_FLAT: {
    int32_t s = (int32_t)read_u32(pos);
    // size in 32 or 64 bit words.
    lbm_uint bind_key = read_u32(pos-1);
    pos -= 2;
    pos -= s;
    lbm_flat_value_t fv;
    fv.buf = (uint8_t*)(image_address + pos);
    fv.buf_size = (uint32_t)s * sizeof(lbm_uint); // GEQ to actual buf
    fv.buf_pos = 0;
    lbm_value unflattened;
    lbm_unflatten_value(&fv, &unflattened);
    if (lbm_is_symbol_merror(unflattened)) {
        lbm_perform_gc();
        lbm_unflatten_value(&fv, &unflattened);
    }

    lbm_uint ix_key = lbm_dec_sym(bind_key) & GLOBAL_ENV_MASK;
    lbm_value *global_env = lbm_get_global_env();
    lbm_uint orig_env = global_env[ix_key];
    lbm_value new_env = lbm_env_set(orig_env,bind_key,unflattened);

    if (lbm_is_symbol(new_env)) {
        return false;
    }
    global_env[ix_key] = new_env;
    pos--;
} break;
```

# Flatten and un-flatten

- Recursive serialisation and de-serialisation functions.
- We ran out of stack on the thread that performed serialisation and de-serialisation.

# Flatten and un-flatten

- Recursive serialisation and de-serialisation functions.
- We ran out of stack on the thread that performed serialisation and de-serialisation.
- Replaced with pointer reversal algorithms!

# Thoughts

- Boot is fast!
- Incremental work against image "possible" but not ideal (duplication).
- Program source only read once, constant block determinism is not a problem.
- A "main" or "startup" function is needed.

# Thoughts

- Flatten and un-flatten works only on non-cyclic lisp values.
- Flatten/un-flatten duplicates shared nodes.
- Don't really need to store the lisp source in flash at all. Future work.

# Duplication

```
(define a (list 1 2 3))
(define b (append (list 'a 'b) a))
(define c (append (list 'c 'd) a))
```

# Conclusion

- Trade-off image size vs performance.
- Trade-off performance vs “correctness”.
- Work around limitations of target platforms.
  - RAM, FLASH, Constant code size concerns.
- Fun.



